Layers and trust boundaries

One palace per process. Tenancy is vaults, not palaces. Nothing leaves the machine unless you configure it to.

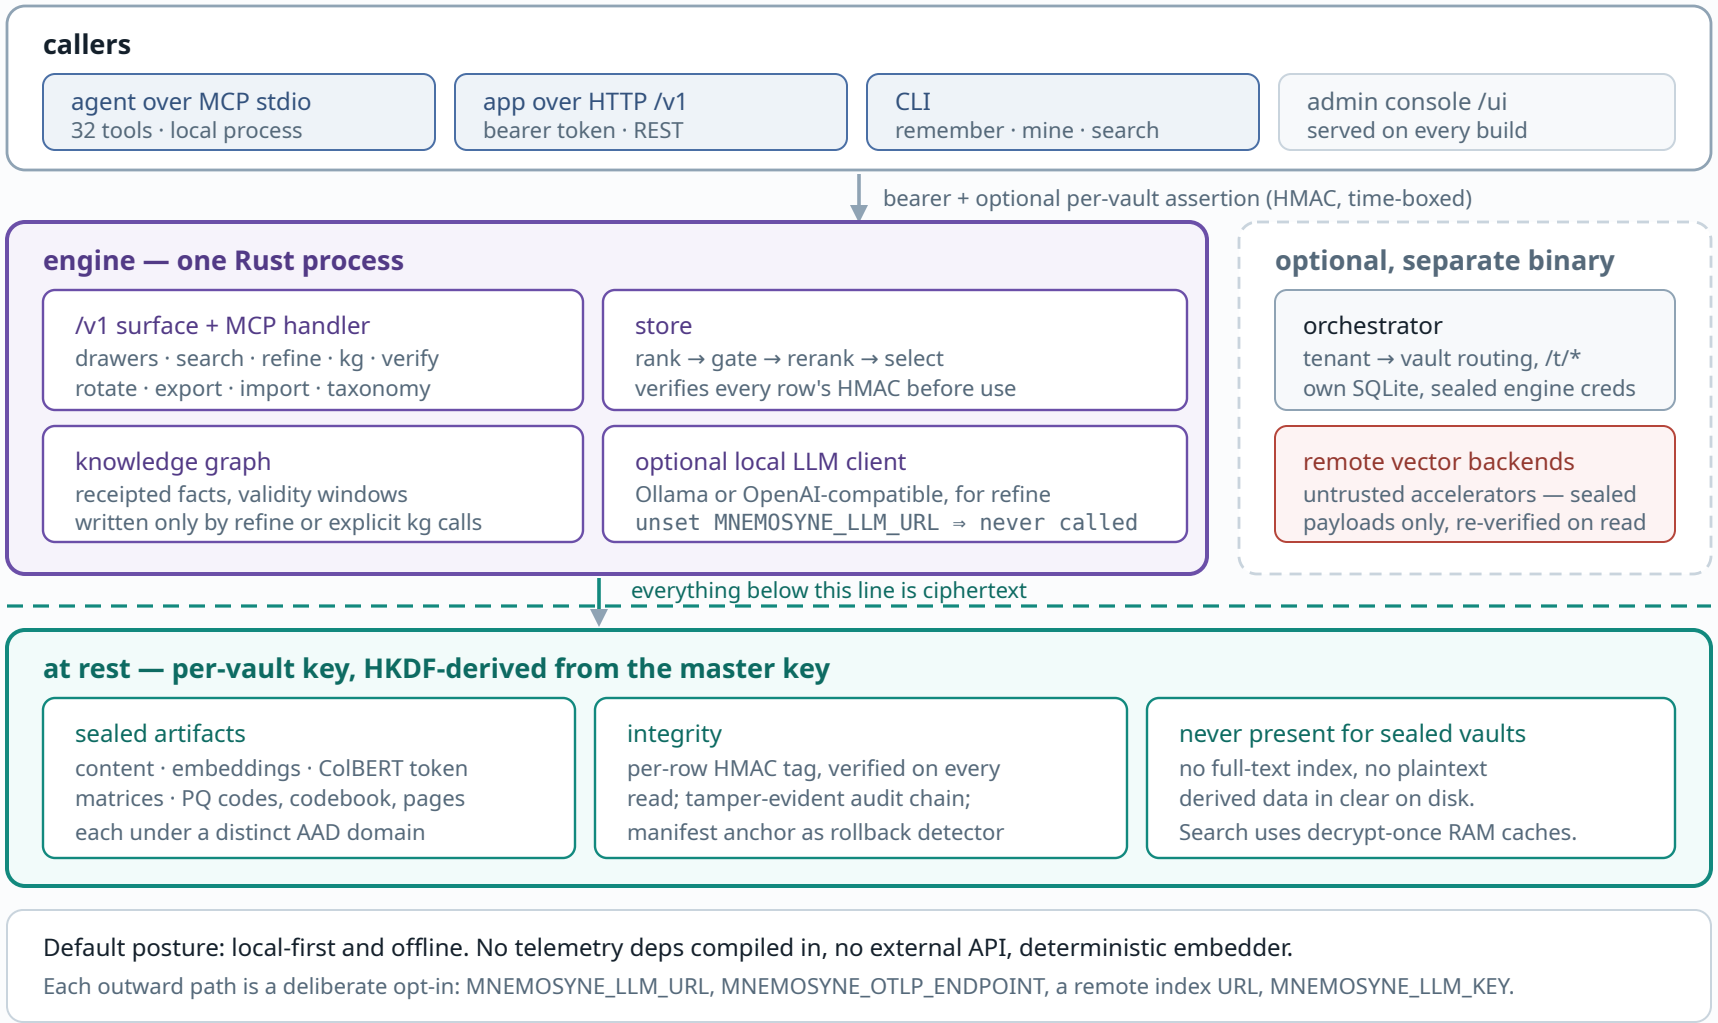

Default posture: local-first and offline. No telemetry deps compiled in, no external API, deterministic embedder.

Each outward path is a deliberate opt-in: MNEMOSYNE_LLM_URL, MNEMOSYNE_OTLP_ENDPOINT, a remote index URL, MNEMOSYNE_LLM_KEY.